

Cloud Run GPU Support

Cloud Run supports GPU-enabled services for machine learning and AI inference workloads.

Key points:

- NVIDIA GPUs can be attached to Cloud Run instances
- GPU services support autoscaling
- GPUs are billed separately from CPU and memory
- GPU instances may have longer cold starts
- Min instances are recommended for latency-sensitive workloads

Supported use cases:

- LLM inference
- Image generation
- Video processing
- Scientific computing

Cloud Run GPU services must still:

- listen on the PORT environment variable
- remain stateless
- use containerized workloads

GPU workloads can connect to:

- Cloud Storage
- Vertex AI
- BigQuery
- Pub/Sub

Cloud Run GPU deployments are configured during deployment using GPU flags in gcloud.